ONLINE MASTERS IN **DATA SCIENCE**

DSC 255 - MACHINE LEARNING FUNDAMENTALS

CROSS-VALIDATION

SANJOY DASGUPTA, PROFESSOR

UC San Diego

COMPUTER SCIENCE & ENGINEERING

HALICIOĞLU DATA SCIENCE INSTITUTE

K-nearest neighbor classification

Classify a point using the labels of its k-nearest neighbors among the training points.

MNIST:

K	1	3	5	7	9	11
Test error (%)	3.09	2.94	3.13	3.10	3.43	3.34

In real life, there's no test set. How to decide which k is best?

K-nearest neighbor classification

Classify a point using the labels of its k-nearest neighbors among the training points.

MNIST:

K	1	3	5	7	9	11
Test error (%)	3.09	2.94	3.13	3.10	3.43	3.34

In real life, there's no test set. How to decide which k is best?

1. **Hold-out set:**
   - Let S be the training set.
   - Choose a subset $V \subset S$ as a validation set.
   - What fraction of V is misclassified by the k-nearest neighbors in $S \setminus V$?

K-nearest neighbor classification

Classify a point using the labels of its k-nearest neighbors among the training points.

MNIST:

K	1	3	5	7	9	11
Test error (%)	3.09	2.94	3.13	3.10	3.43	3.34

In real life, there's no test set. How to decide which k is best?

1. **Hold-out set:**
   - Let S be the training set.
   - Choose a subset $V \subset S$ as a *validation* set.
   - What fraction of V is misclassified by the *k*-nearest neighbors in $S \setminus V$?
2. **Leave-one-out cross-validation:**
   - For each point $x \in S$, find the *k*-nearest neighbors in $S \setminus \{x\}$.
   - What fraction are misclassified?

Cross-validation

How to estimate the error of *k*-NN for a particular *k*?

10-fold cross-validation

- Divide the training set into 10 equal pieces.
- Training set (call it S): 60,000 points
- Call the pieces S_1, S_2, …, S_{10}: 6,000 points each.

- For each piece S_i:
  - Classify each point in S_i using *k*-NN with training set $S - S_i$
  - Let ϵ_i = fraction of S_i that is incorrectly classified

- Take the average of these 10 numbers:

$$\text{estimated error with } k\text{-NN} = \frac{\epsilon_1 + \cdots + \epsilon_{10}}{10}$$