

# Module 5: Optimization and Gradient Descent

Machine Learning Course

## Contents

<b>1</b>	<b>Introduction to Optimization in Machine Learning</b>	<b>3</b>
1.1	The Optimization Framework . . . . .	3
1.2	Common Loss Functions . . . . .	3
1.3	Local Search Methods . . . . .	3
<b>2</b>	<b>Convexity</b>	<b>3</b>
2.1	Definition and Intuition . . . . .	3
2.2	Why Convexity Matters in Optimization . . . . .	4
2.3	Checking Convexity . . . . .	4
2.4	Positive Semidefinite (PSD) Matrices . . . . .	4
2.5	Worked Examples . . . . .	5
<b>3</b>	<b>Multivariate Differentiation</b>	<b>6</b>
<b>4</b>	<b>Gradient Descent</b>	<b>6</b>
4.1	The Gradient and Its Geometric Meaning . . . . .	6
4.2	Gradient Descent Algorithm . . . . .	7
4.3	Rationale: Local Linearity and Descent Direction . . . . .	7
4.4	How to Set Step Size $\eta_t$ ? . . . .	7
4.5	Gradient Descent for Logistic Regression: Full Derivation . . . . .	7
<b>5</b>	<b>Variants of Gradient Descent</b>	<b>8</b>
5.1	Decomposable Loss Functions . . . . .	8
5.2	Stochastic Gradient Descent (SGD) . . . . .	8
5.3	Mini-Batch Gradient Descent . . . . .	8
5.4	Comparison Table: GD, SGD, Mini-Batch GD . . . . .	8
<b>6</b>	<b>Convergence Properties</b>	<b>8</b>
6.1	Convergence Guarantees for Convex Functions . . . . .	8
6.2	Rates for GD, SGD, and Accelerated Methods . . . . .	9
6.3	Practical Convergence Criteria (Stopping Conditions) . . . . .	9
<b>7</b>	<b>Advanced Optimization Methods</b>	<b>9</b>
7.1	Momentum . . . . .	9
7.2	Adaptive Learning Rates . . . . .	9
7.3	Second-Order Methods . . . . .	9

<b>8</b>	<b>Practical Considerations</b>	<b>10</b>
8.1	Initialization Strategies . . . . .	10
8.2	Regularization and Its Effect . . . . .	10
8.3	Non-Convex Optimization: Challenges and Context . . . . .	10
8.4	Tips for Debugging and Tuning Optimization . . . . .	10
<b>9</b>	<b>Practical Considerations</b>	<b>10</b>
9.1	Initialization Strategies . . . . .	10
9.2	Regularization and Its Effect . . . . .	10
9.3	Non-Convex Optimization: Challenges and Context . . . . .	10
9.4	Tips for Debugging and Tuning Optimization . . . . .	11

# 1 Introduction to Optimization in Machine Learning

## 1.1 The Optimization Framework

In machine learning, we typically choose a model parameterized by  $w$  by minimizing a loss function  $L(w)$  that depends on the training data. This optimization problem is central to most machine learning algorithms.

## 1.2 Common Loss Functions

Different machine learning tasks use different loss functions:

- **Linear Regression:**

$$L(w) = \sum_{i=1}^n \left( y^{(i)} - w \cdot x^{(i)} \right)^2$$

- **Logistic Regression:**

$$L(w) = \sum_{i=1}^n \ln \left( 1 + e^{-y^{(i)}(w \cdot x^{(i)})} \right)$$

where  $y^{(i)} \in \{-1, 1\}$ .

## 1.3 Local Search Methods

The default approach to solving these minimization problems is through local search:

1. Initialize  $w$  arbitrarily
2. Repeat until convergence:
  - (a) Find some  $w'$  close to  $w$  with  $L(w') < L(w)$
  - (b) Move  $w$  to  $w'$

Local search methods work particularly well when the loss function is convex, which we'll explore in detail later in this module.

# 2 Convexity


## 2.1 Definition and Intuition

A function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  is convex if for all  $a, b \in \mathbb{R}^d$  and  $0 < \theta < 1$ :

$$f(\theta a + (1 - \theta)b) \leq \theta f(a) + (1 - \theta)f(b)$$

Intuitively, this means that the line segment connecting any two points on the graph of the function lies above or on the graph.

A function is strictly convex if strict inequality holds for all  $a \neq b$ . Conversely,  $f$  is concave if  $-f$  is convex.



Convex Function

Figure 1: Illustration of convexity: The function value at any point on the line segment connecting two points is less than or equal to the weighted average of the function values at those points.

## 2.2 Why Convexity Matters in Optimization

Convexity ensures that any local minimum is a global minimum. This property is crucial for the success of local search and gradient-based optimization methods.

## 2.3 Checking Convexity

### One Variable: Second Derivative Test

A twice-differentiable function  $f : \mathbb{R} \rightarrow \mathbb{R}$  is convex if and only if  $f''(x) \geq 0$  for all  $x$  in the domain.

### Multivariate: Hessian and PSD Matrices

For  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ , the Hessian matrix  $H(x)$  is defined as:

$$H_{jk}(x) = \frac{\partial^2 f}{\partial x_j \partial x_k}(x)$$

A twice-differentiable function is convex if and only if its Hessian is positive semidefinite (PSD) everywhere.

## 2.4 Positive Semidefinite (PSD) Matrices

A symmetric matrix  $M$  is PSD if  $x^T M x \geq 0$  for all  $x \in \mathbb{R}^d$ .

### Properties:

- Diagonal matrix is PSD if all diagonal entries  $\geq 0$ .
- If  $M$  is PSD and  $c > 0$ , then  $cM$  is PSD.
- If  $M, N$  are PSD, then  $M + N$  is PSD.

- $M$  is PSD iff  $M = UU^T$  for some  $U$ .
- All covariance matrices are PSD.

## 2.5 Worked Examples

**Example 1: Convexity of  $f(x) = \|x\|^2$**

$$f(x) = \|x\|^2 = \sum_{i=1}^d x_i^2$$

$$\frac{\partial f}{\partial x_j} = 2x_j$$

$$\frac{\partial^2 f}{\partial x_j \partial x_k} = 2\delta_{jk}$$

So the Hessian is  $2I$ , which is positive definite. Therefore,  $f(x)$  is strictly convex.

**Example 2: Convexity of  $f(z) = (u \cdot z)^2$**

$$f(z) = (u \cdot z)^2$$

$$\frac{\partial f}{\partial z_j} = 2(u \cdot z)u_j$$

$$\frac{\partial^2 f}{\partial z_j \partial z_k} = 2u_j u_k$$

So the Hessian is  $2uu^T$ , which is PSD since for any  $x$ ,  $x^T(2uu^T)x = 2(u^T x)^2 \geq 0$ .

**Example 3: Is  $M = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$  PSD?** Let  $x = (x_1, x_2)^T$ :

$$x^T M x = (x_1 + x_2)^2 \geq 0$$

So  $M$  is PSD.

**Example 4: Is  $M = \begin{pmatrix} 1 & 2 \\ 2 & 1 \end{pmatrix}$  PSD?** Eigenvalues are 3 and  $-1$ . Since one eigenvalue is negative,  $M$  is not PSD.

**Example 5: Convexity of Least Squares Loss**

$$L(w) = \sum_{i=1}^n (y^{(i)} - w \cdot x^{(i)})^2$$

The Hessian is  $2 \sum_{i=1}^n x^{(i)}(x^{(i)})^T$ , a sum of PSD matrices, so  $L$  is convex.

### Example 6: Convexity of Logistic Regression Loss

$$L(w) = \sum_{i=1}^n \ln(1 + e^{-y^{(i)}(w \cdot x^{(i)})})$$

Each term  $\ell(z) = \ln(1 + e^{-z})$  has  $\ell''(z) > 0$ , so  $L$  is convex.

## 3 Multivariate Differentiation

For  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ :

- **Gradient:**  $\nabla f(w) = \left( \frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T$
- **Hessian:**  $H_{jk}(w) = \frac{\partial^2 f}{\partial w_j \partial w_k}$

**Example 1:**  $F(w) = 3w_1w_2 + w_3$  for  $w \in \mathbb{R}^3$

$$\frac{\partial F}{\partial w_1} = 3w_2, \quad \frac{\partial F}{\partial w_2} = 3w_1, \quad \frac{\partial F}{\partial w_3} = 1$$

So,

$$\nabla F(w) = (3w_2, 3w_1, 1)^T$$

**Example 2:**  $F(w) = w \cdot x$  where  $x$  is fixed

$$\frac{\partial F}{\partial w_j} = x_j$$

So,

$$\nabla F(w) = x$$

**Example 3:** Second Derivative Matrix of  $F(w) = \|w\|^2$

$$F(w) = \sum_{j=1}^d w_j^2$$
$$\frac{\partial^2 F}{\partial w_j \partial w_k} = 2\delta_{jk}$$

So the Hessian is  $2I$ .

## 4 Gradient Descent

### 4.1 The Gradient and Its Geometric Meaning

The gradient points in the direction of steepest increase of a function. Moving in the negative gradient direction decreases the function most rapidly.

## 4.2 Gradient Descent Algorithm

1. Initialize  $w_0 = 0, t = 0$
2. While  $\|\nabla L(w_t)\| > \epsilon$  (not converged):
  - (a)  $w_{t+1} = w_t - \eta_t \nabla L(w_t)$
  - (b)  $t = t + 1$

Here,  $\eta_t > 0$  is the step size (learning rate).

## 4.3 Rationale: Local Linearity and Descent Direction

For small  $u$ ,

$$L(w + u) \approx L(w) + u \cdot \nabla L(w)$$

Choosing  $u = -\eta \nabla L(w)$  for small  $\eta$  ensures  $L(w + u) < L(w)$ .

## 4.4 How to Set Step Size $\eta_t$ ?

- **Constant step size:**  $\eta_t = \eta$
- **Diminishing step size:**  $\eta_t = \frac{\eta_0}{1+\beta t}$  or  $\eta_t = \frac{\eta_0}{\sqrt{t}}$
- **Backtracking line search:** Iteratively decrease  $\eta_t$  until  $L(w_t - \eta_t \nabla L(w_t)) < L(w_t)$
- **Adaptive methods:** Adjust  $\eta_t$  based on gradient history (e.g., AdaGrad, Adam)

## 4.5 Gradient Descent for Logistic Regression: Full Derivation

Given  $(x^{(i)}, y^{(i)}) \in \mathbb{R}^d \times \{-1, 1\}$ ,

$$L(w) = \sum_{i=1}^n \ln(1 + e^{-y^{(i)}(w \cdot x^{(i)})})$$

$$\nabla L(w) = \nabla \sum_{i=1}^n \ln(1 + e^{-y^{(i)}(w \cdot x^{(i)})}) \quad (1)$$

$$= \sum_{i=1}^n \nabla \ln(1 + e^{-y^{(i)}(w \cdot x^{(i)})}) \quad (2)$$

$$= \sum_{i=1}^n \frac{1}{1 + e^{-y^{(i)}(w \cdot x^{(i)})}} \cdot \nabla e^{-y^{(i)}(w \cdot x^{(i)})} \quad (3)$$

$$= \sum_{i=1}^n \frac{1}{1 + e^{-y^{(i)}(w \cdot x^{(i)})}} \cdot e^{-y^{(i)}(w \cdot x^{(i)})} \cdot (-y^{(i)} x^{(i)}) \quad (4)$$

$$= - \sum_{i=1}^n \frac{e^{-y^{(i)}(w \cdot x^{(i)})}}{1 + e^{-y^{(i)}(w \cdot x^{(i)})}} \cdot y^{(i)} x^{(i)} \quad (5)$$

$$= - \sum_{i=1}^n \frac{1}{1 + e^{y^{(i)}(w \cdot x^{(i)})}} \cdot y^{(i)} x^{(i)} \quad (6)$$

$$= - \sum_{i=1}^n P(Y \neq y^{(i)} | x^{(i)}, w) \cdot y^{(i)} x^{(i)} \quad (7)$$

Update rule:

$$w_{t+1} = w_t - \eta_t \nabla L(w_t)$$

## 5 Variants of Gradient Descent

### 5.1 Decomposable Loss Functions

Many loss functions decompose as  $L(w) = \sum_{i=1}^n \ell(w; x^{(i)}, y^{(i)})$ .

### 5.2 Stochastic Gradient Descent (SGD)

Update using a single data point:

$$w_{t+1} = w_t - \eta_t \nabla \ell(w_t; x^{(i)}, y^{(i)})$$

### 5.3 Mini-Batch Gradient Descent

Update using a batch  $B$ :

$$w_{t+1} = w_t - \eta_t \sum_{(x,y) \in B} \nabla \ell(w_t; x, y)$$

### 5.4 Comparison Table: GD, SGD, Mini-Batch GD

## 6 Convergence Properties

### 6.1 Convergence Guarantees for Convex Functions

For convex  $L(w)$  and appropriate step sizes, gradient descent converges to the global minimum.



Method	Update Rule	Advantages	Disadvantages
Gradient Descent	$w_{t+1} = w_t - \eta_t \nabla L(w_t)$	Stable, accurate	Slow for large d
SGD	$w_{t+1} = w_t - \eta_t \nabla \ell(w_t; x^{(i)}, y^{(i)})$	Fast, low memory, escapes local minima	Noisy, may not
Mini-Batch GD	$w_{t+1} = w_t - \eta_t \sum_{(x,y) \in B} \nabla \ell(w_t; x, y)$	Balanced, parallelizable	Needs batch size

Table 1: Comparison of gradient descent variants

## 6.2 Rates for GD, SGD, and Accelerated Methods

- GD:  $O(1/T)$  for general convex, linear for strongly convex.
- SGD:  $O(1/\sqrt{T})$  in expectation.
- Accelerated (e.g., Nesterov):  $O(1/T^2)$  for smooth convex.

## 6.3 Practical Convergence Criteria (Stopping Conditions)

- $\|\nabla L(w_t)\| < \epsilon$
- $|L(w_{t+1}) - L(w_t)| < \delta$
- Maximum number of iterations

# 7 Advanced Optimization Methods

## 7.1 Momentum

$$v_{t+1} = \gamma v_t + \eta \nabla L(w_t)$$

$$w_{t+1} = w_t - v_{t+1}$$

where  $\gamma \in [0, 1)$ .

## 7.2 Adaptive Learning Rates

- AdaGrad: Per-parameter learning rates, decreases for frequent features.
- RMSProp: Exponential moving average of squared gradients.
- Adam: Combines momentum and RMSProp.

## 7.3 Second-Order Methods

- Newton's Method:  $w_{t+1} = w_t - [H(w_t)]^{-1} \nabla L(w_t)$
- Quasi-Newton (e.g., BFGS, L-BFGS): Approximate inverse Hessian.

## 8 Practical Considerations

### 8.1 Initialization Strategies

- For convex problems, any  $w_0$  will converge.
- For non-convex (e.g., neural nets), initialization affects which minimum is found.
- Common: random, Xavier/Glorot, He initialization.

### 8.2 Regularization and Its Effect

Adding regularization (e.g., L2:  $\lambda\|w\|^2$ ) can improve convexity and generalization.

### 8.3 Non-Convex Optimization: Challenges and Context

Many modern ML models (e.g., deep neural networks) are non-convex. While gradient descent can still work well in practice, there are no guarantees of finding the global minimum.

### 8.4 Tips for Debugging and Tuning Optimization

- Monitor loss and gradients.
- Try different learning rates and batch sizes.
- Use validation data to check for overfitting.
- Visualize convergence when possible.

## 9 Practical Considerations

### 9.1 Initialization Strategies

- For convex problems, any  $w_0$  will converge.
- For non-convex (e.g., neural nets), initialization affects which minimum is found.
- Common: random, Xavier/Glorot, He initialization.

### 9.2 Regularization and Its Effect

Adding regularization (e.g., L2:  $\lambda\|w\|^2$ ) can improve convexity and generalization.

### 9.3 Non-Convex Optimization: Challenges and Context

Many modern ML models (e.g., deep neural networks) are non-convex. While gradient descent can still work well in practice, there are no guarantees of finding the global minimum.

## 9.4 Tips for Debugging and Tuning Optimization

- Monitor loss and gradients.
- Try different learning rates and batch sizes.
- Use validation data to check for overfitting.
- Visualize convergence when possible.