

## Week 9 — Solutions

1. *Other methods with the expressive power of decision trees.*

- (a) Linear classifiers cannot represent any boundary that is not linear.
- (b) Support vector machines with a quadratic kernel cannot represent any boundary that is not quadratic.
- (c) Nearest neighbor can capture any boundary.
- (d) Classifiers based on Gaussian generative models end up with quadratic boundaries and are thus not as expressive as decision trees.

2. There are  $d$  features, and  $n - 1$  split points for each feature (look at the one-dimensional values of that feature alone, in sorted order; the split will lie between two consecutive values). Thus the total number of possible splits is  $d(n - 1)$ .3. When  $p = 0.2$ , the Gini index is  $2 \times 0.2 \times 0.8 = 0.32$ .4. *Working with weighted data.* Suppose we have data points  $(x_1, y_1), \dots, (x_n, y_n)$ , with weights  $\lambda_1, \dots, \lambda_n > 0$ . Suppose the possible labels are  $\{1, 2, \dots, k\}$ .

- (a) At any node of the data, let  $S \subset [n]$  denote the subset of points reaching that node. We compute the proportions of the different labels  $(p_1, \dots, p_k)$  using the weights of these points:

$$p_j = \frac{\sum_{i \in S} \lambda_i \cdot 1(y_i = j)}{\sum_{i \in S} \lambda_i}.$$

The impurity of a split (e.g., using the Gini index) is then a function of these  $p_j$  values.

- (b) To compute the mean and covariance for class  $j$ , use the weights:

$$\mu_j = \frac{\sum_{i=1}^n \lambda_i \cdot 1(y_i = j) x_i}{\sum_{i=1}^n \lambda_i \cdot 1(y_i = j)}, \quad \Sigma_j = \frac{\sum_{i=1}^n \lambda_i \cdot 1(y_i = j) (x_i - \mu_j)(x_i - \mu_j)^T}{\sum_{i=1}^n \lambda_i \cdot 1(y_i = j)}.$$

- (c) Include weights in the SVM optimization: e.g., for the binary case, use

$$\begin{aligned} \min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad & \|w\|^2 + C \sum_{i=1}^n \lambda_i \xi_i \\ \text{s.t.:} \quad & y_i(w \cdot x_i + b) \geq 1 - \xi_i \quad \text{for all } i = 1, 2, \dots, n \\ & \xi \geq 0 \end{aligned}$$

5. *Convergence behavior of boosting.*

- (a) There is no guarantee that boosting will converge to a model with zero test error. Such a model need not even exist, e.g., in situations with inherent uncertainty (recall our discussion of cases where perfect classification might not be possible).

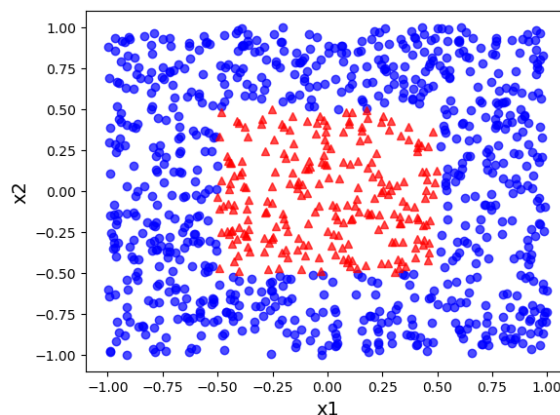
- (b) True, boosting will converge to a model with zero training error.
- (c) False, because boosting does not converge to a model in  $H$ . Its model is a linear combination of classifiers from  $H$ .

6. *Random forests versus boosted decision trees.*

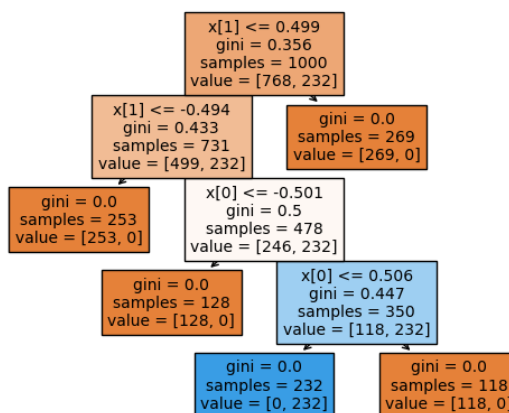
- (a) True, random forests can be trained in parallel whereas in boosted decision trees, the trees must be trained sequentially.
- (b) False: each individual tree in a random forest is not more highly optimized.
- (c) False: each individual tree in a random forest need not have better accuracy.

7. *A toy 2-d data set for decision trees and boosting.*

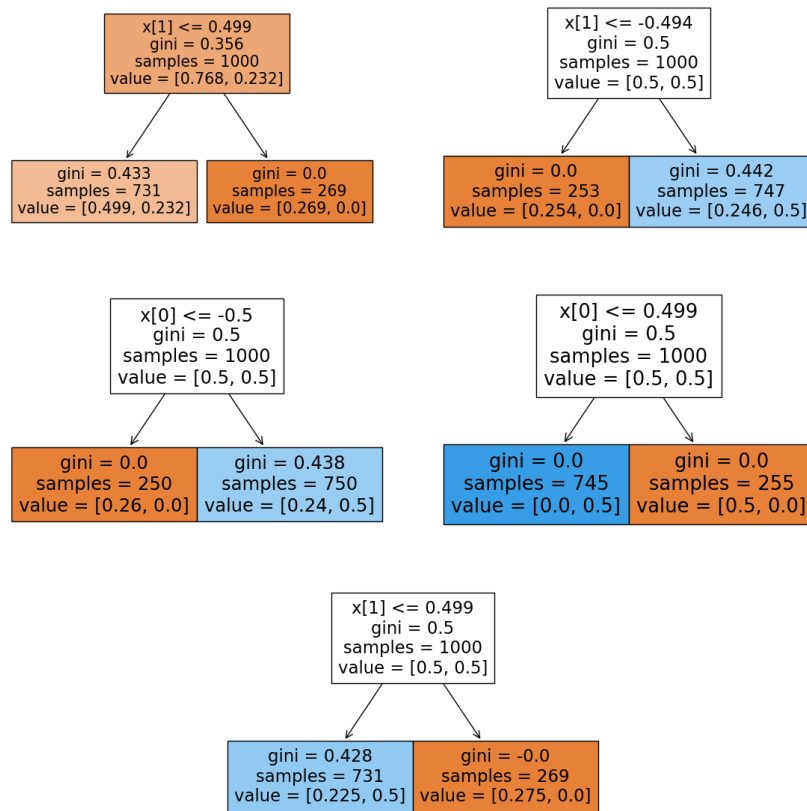
- (a) Here is the data from `mini-data.txt`.



- (b) There are many reasonable stopping criteria. We used `max_depth = 4`.
- (c) Here is the decision tree.



- (d) We set the boosting iterations to 5. Here are the stumps.

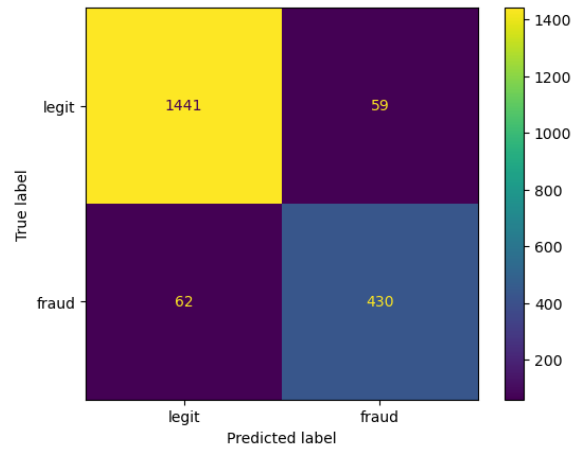


(e) Here is a table of training accuracy with each successive stump.

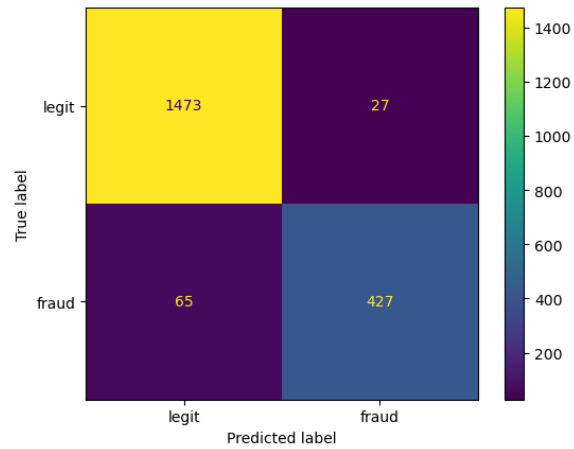
# Stumps	Accuracy
1	0.768
2	0.768
3	0.882
4	1.000
5	1.000

8. *Credit card fraud data.*

- Out of the 284,807, only 492 are fraudulent. This is problematic because a classifier can get low error by always predicting legitimate.
- We retain all 492 fraudulent transactions and subsample just 1500 legitimate transactions. This gives a training set of size 1992.
- Here are the confusion matrices for the three methods.  
Decision tree:



Boosted decision stumps:



Random forest:

